

A Method for Regularisation of Evolutionary Polynomial Regressions using Support Vector Machines

Francisco Coelho^{a,c,*}, João Pedro Neto^{b,c}

^a*Dept. Informática, Universidade de Évora, Rua Romão Ramalho 58, 7000-671 Évora*

^b*Dept. Informática, Faculdade de Ciências da Universidade de Lisboa, Campo Grande
1749-016 Lisboa*

^c*Laboratory of Agent Modelling (LabMAg)*

Abstract

Many applications require models that have no acceptable linear approximation and the simpler nonlinear models are defined by polynomials. The use of genetic algorithms to find polynomial models, Evolutionary Polynomial Regression (EPR), provides an automatic method to build such model from data but still poses challenges due to the complexity of the search and different definitions of optimal solution.

This paper describes a two-step method that uses genetic algorithms and linear regression to find empirical polynomial regressions. Experiments on common datasets show that, discounting the training computational effort, this method is quite competitive.

Keywords: evolutionary polynomial regression, regularization, feature extraction, dimensionality reduction

*Corresponding author

Email addresses: `fc@di.uevora.pt` (Francisco Coelho), `jpn@di.fc.ul.pt` (João Pedro Neto)

1. Introduction

With notable exceptions (*e.g.* neural networks) machine learning regression techniques are based on linear models. The linearity assumption has many advantages including reduced computational complexity and strong theoretical framework. However nonlinearity is unavoidable in many application scenarios, specially those with phase transitions or feedback loops, so common in engineering, ecology, cybernetics and other areas. The kernel trick in Support Vector Machines alleviates this problem by allowing a non-linear transformation of the feature-space. [**Find references.**]

Polynomials, one of the most studied subjects in mathematics, generalize linear functions and define, perhaps, the simplest and most used nonlinear models. Applications include colorimetric calibration (Mendes and Carvalho, 2005), explicit formulæ for turbulent pipe flows (Davidson et al., 1999), computational linguistics (Sánchez et al., 2009) and more recently, analytical techniques for cultural heritage materials (Cséfalvayová et al., 2010), liquid epoxy molding process (Chan et al., 2011), B-spline surface reconstruction (Gálvez et al., 2012), product design (Chan et al., 2012) or forecasting cyanotoxins presence in water reservoirs (García Nieto et al., 2013). These examples not only illustrate the wide spectrum of applications but, additionally, work in each one uses, at some point, a genetic algorithm.

Genetic algorithms (GA) where, arguably, one of the hottest topics of research in the recent decades and with good reason since they outline an optimization scheme easy to conceptualize and with very broad application. If a nonlinear (or otherwise) model requires parameterization GAs provide a simple and often effective approach to search for locally optimal parameters.

26 Research related to genetic algorithms abound and spans from the 1950s
 27 seminal work of Nils Aall Barricelli (Barricelli, 1962) in the Institute for Ad-
 28 vanced Study of Princeton to today’s principal area of study for thousands of
 29 researchers, covered in hundreds of conferences, workshops and other meet-
 30 ings. Perhaps the key impulse to GAs come from John Holland’s work and
 31 his book “Adaptation in Natural and Artificial Systems” (Holland, 1975).

32 One interesting take on genetic algorithms, named *genetic programming*
 33 by John Koza (Koza, 1992), proposed the use of GAs to search the syntactic
 34 structure of complex functions. This syntactic structure search is keen to the
 35 central ideas of deep learning (Bengio et al., 2013; Bengio, 2009), a subarea of
 36 machine learning actually producing quite promising results (*e.g.* in Tarlow
 37 et al. (2013)). It is also related to the work presented in this paper in the
 38 sense that, unlike linear models that have a simple structure, $y = \sum_i \beta_i x_i$,
 39 nonlinear (in particular polynomial) models pose an additional “structure”
 40 search problem.

41 The idea of using GAs to find a polynomial regression is not new (Maertens
 42 et al., 2006; Yu and Lin, 2008; Wu et al., 2009) but still generates original
 43 research (Hofwing et al., 2011; Cetisli and Kalkan, 2011). [**Insert refer-**
 44 **ences and resume of EPR; Redo the purpose of this work.**] In
 45 line with that research this work describes a general method to find a poly-
 46 nomial regression of a given dataset. The optimal regression minimizes a
 47 cost function that accounts for both the root-mean-square error (error) and
 48 a regularization factor to avoid over-fitting.

49 It turns out that, discarding the computational cost of training, the poly-
 50 nomial regression method presented here, Genetic Algorithms for Polynomi-

als (GAPOLY), provides a quite competitive regression method. Indeed, it is only systematically out-performed by random forests, an *ensemble* method.

The remainder of this paper is organized as usual: the next section describes the details of our method and is followed by a presentation of some performance results. The last section draws some conclusions and points future research tasks.

2. Genetic Algorithms for Polynomials

This section is dedicated to the description of an algorithm to find a polynomial regression from a given dataset. It starts with a brief introduction and outline of the algorithm and proceeds into core details as the encoding used to represent individual polynomial instances in the GA populations and the regularization of the cost function.

An usual representation of polynomials is

$$p(x_1, \dots, x_m) = \sum_i \theta_i q_i$$

where each q_i is a monomial, $q_i = \prod_j x_j^{\alpha_{ij}}$, the exponents are non-negative integers, $\alpha_{ij} \in \mathbb{N}_0$, and the coefficients are real valued, $\theta_i \in \mathbb{R}$. For example $p(x_1, x_2, x_3) = 2x_1 + x_2x_3 + \frac{1}{2}x_1^2x_3$ has monomials $q_1 = x_1$, $q_2 = x_2x_3$ and $q_3 = x_1^2x_3$, coefficients $\theta_1 = 2$, $\theta_2 = 1$ and $\theta_3 = 1/2$ and exponents $\alpha_{1,1} = 1$, $\alpha_{2,2} = 1$, $\alpha_{2,3} = 1$, $\alpha_{3,1} = 2$, $\alpha_{3,3} = 1$ and all other $\alpha_{ij} = 0$. The exponents alone are a matrix that defines the monomial structure of the polynomial, $A = [\alpha_{ij}]$.

For the example above the matrix of monomials is

$$A = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 1 \\ 2 & 0 & 1 \end{bmatrix} \sim \begin{bmatrix} x_1 \\ x_2 x_3 \\ x_1^2 x_3 \end{bmatrix}$$

70 where each row defines a monomial and each column represents a variable.

71 Changing the order of the rows doesn't change the polynomial whereas chang-

72 ing the order of the columns corresponds to changing the respective variables.

73 This representation of polynomials makes the problem of structure search

74 very clear: except for the trivial cases, the number of possible monomials

75 given n variables and a maximum joint degree d grows exponentially with

76 either n or d . But more importantly, the polynomial regression problem can

77 be naturally split into two subproblems:

78 1. For a given set of monomials $\mathcal{P} = \{q_1, \dots, q_k\}$, find the regression
79 coefficients $\theta_1, \dots, \theta_k$ that minimize the error on a given dataset;

80 2. Find the fittest set of monomials, *i.e.* the polynomial that minimizes
81 the error on the same dataset;

82 More precisely, concerning the first problem, let \mathcal{D} be a dataset with n obser-

83 vations of variables Y, X_1, \dots, X_m and $\mathcal{P} = \{q_1, \dots, q_k\}$ a set of k monomial

84 expressions over X_1, \dots, X_m . Define the hypothesis¹

$$h_{\Theta, \mathcal{P}}(x_1, \dots, x_m) = \sum_{j=1}^k \theta_j q_j|_{X_i=x_i, \forall 1 \leq i \leq m}$$

85 and let the cost

$$J_{\text{fit}}(\Theta; \mathcal{P}, \mathcal{D}) =$$

¹The expression $q|_{X=x}$ reads “replace all instances of X by x in q ”.

$$\sqrt{\frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - h_{\Theta, \mathcal{P}} \left(x_1^{(i)}, \dots, x_m^{(i)} \right) \right)^2} \quad (1)$$

86 be the usual root-mean-square error (error) function. Now the first prob-
 87 lem can be stated as: *Given a dataset \mathcal{D} and a set of monomials \mathcal{P} , find*
 88 *parameters Θ that minimize $J_{\text{fit}}(\Theta; \mathcal{P}, \mathcal{D})$.*

89 It turns out that this problem can be solved as a usual linear regression
 90 problem by expanding \mathcal{D} with columns that replicate the monomials in \mathcal{M} .

91 The second problem is treated in the GA setting: Let \mathcal{D} be a dataset as
 92 above and Q a set of polynomials. For each polynomial $p \in Q$ let \mathcal{P}_p be the
 93 set of monomials in p (without the coefficients) and compute the fitness

$$\phi_p = \min_{\Theta} J_{\text{fit}}(\Theta; \mathcal{P}_p, \mathcal{D})$$

94 by solving the first problem. With a fitness of every instance, a GA will apply
 95 genetic operators (usually mutation and crossover) to evolve the population
 96 Q until a reasonable approximation of a local minimum is found. Notice
 97 that the properties of GAs and linear regression entail that the composition
 98 of GAs with linear regression, as defined in Algorithm 1, converges to a
 99 polynomial that is a local minimum of the fitness function, encapsulated in
 100 the error function J_{fit} .

101 Subsection 2.1 describes the encoding of individual polynomial instances
 102 as chromosomes and other parameters of the utilized GA implementation.
 103 The regularization of the cost function is discussed in subsection 2.2.

104 2.1. Polynomial Encoding

105 The encoding for any polynomial will be as follows:

Algorithm 1 GAPOLY uses linear regression to find monomial coefficients that minimize the error over a dataset and GAs to explore the space of polynomials. The role of the linear regression step is to produce a fitting error that sorts the population. At exit the error of the fittest instance is bounded by ϵ .

```

function GAPOLY( $D, pop_0, \epsilon$ )
     $pop \leftarrow pop_0; err \leftarrow 1.0 + \epsilon$ 
    while  $err > \epsilon$  do
         $pop \leftarrow \text{ITERATEGA}(pop)$ 
         $pop \leftarrow \text{SORT}(pop, key = J)$   $\triangleright$  Sort population by regression error
         $err \leftarrow J(\text{FIRST}(pop))$ 
    end while
    return  $\text{FIRST}(pop)$ 
end function

```

- 106 1. an initial segment detailing which monomials are active (the first mono-
- 107 mial is always active), this is represented in unary description, i.e., each
- 108 monomial is identified by a single bit
- 109 2. the remaining bits are split into k sets of equal size, each one repre-
- 110 senting a monomial
- 111 3. each monomial is split into m sets of d size each, i.e., a variable
- 112 4. for each variable, the remaining bits are the binary description of the
- 113 variable degree, i.e., the maximum exponent is given by $2^d - 1$
- 114 Let's see an example: consider polynomial $x_1^3x_3 + x_3^7 + x_1x_2$ with $k =$
- 115 $4, m = 3$ and $d = 3$. One possible encoding would be:

116 110 - 011,000,001 ; 000,000,111 ; 001,001,000 ; 110,010,101

117 (for reading purposes the semicolons separate monomials, the commas
118 separate variables)

119 The first three bits inform that the second and third monomials are active
120 while the fourth is not (as said, the first monomial is always active). This
121 last monomial does not enter neither in the polynomial regression nor in the
122 error (fitness) evaluation. However, it acts as a kind of junk DNA, becoming
123 active when, in a future mutation or crossover, the third bit of the entire
124 sequence flips from 0 to 1.

125 Let's interpret the first monomial description, 011,000,001. It is divided
126 by three since $m = 3$. The first triple 011 is the binary description of the
127 exponent of variable x_1 which is 3, so the first monomial includes x_1^3 . The
128 second triple, 000, means that x_2 is not part of the monomial. The third
129 triple 001 says that variable x_3 has exponent 1. The first monomial is $x_1^3x_3$.

130 Notice that all binary descriptions give rise to valid polynomials. However
131 this is not a bijective mapping. For each polynomial there are multiple rep-
132 resentations. For example, $x_1 + x_2$ and $x_2 + x_1$ have different representations.
133 The authors considered that more complex mappings in order to force a one
134 to one mapping would impact negatively in the algorithm's performance.

135 2.2. Cost Function

136 The polynomial regression error considered so far is based on the ability to
137 predict the test set after the polynomial regression has found the appropriate
138 coefficients θ_i for each one of the monomials q_i .

139 This error function tends to prefer more complex polynomials, namely in
140 the number of monomials which provides the regression algorithm for more




Figure 1: Error distribution by regularization exponent for the Abalone and Auto MPG datasets. The box plots summarize the error values of 10 runs for each value of λ . The smallest overall error is shown in gray and, in both cases, was achieved with $\lambda = 0.8$.

141 fitting possibilities. One way to balance this is to provide a regularization
 142 term into the error function. Our proposal is to include a multiplicative factor
 143 proportional to the number of monomials. Thus, the error from equation 1
 144 defines

$$J_{\text{reg}}(\Theta; \mathcal{P}, \mathcal{D}) = \lambda^k J_{\text{fit}}(\Theta; \mathcal{P}, \mathcal{D}) \quad (2)$$

145 where k is the number of monomials in the polynomial. When $\lambda > 1$ poly-
 146 nomials with more monomials are penalized.

147 Somewhat unexpectedly after some experiences it was found that lower
 148 values for λ provide better, even if marginal, results. Figure 1 shows regu-
 149 larization results for the Abalone and Auto MPG datasets with ten runs for
 150 each λ . The following section includes information about these datasets.

151 The typical inflection point lies around $\lambda = 0.8$. The dataset results for
 152 applying the proposed regression method use the regularization parameter
 153 with this value.

154 2.3. Genetic Operators

155 To perform the genetic algorithm it was used the R package `genalg` (Wil-
156 lighagen, 2012). The operators were the standard ones: (a) crossover, *i.e.*,
157 a pair of solutions from the previous generations are combined by splitting
158 and mixing their respective representations, and (b) mutation, changing the
159 values of single bits; the mutation chance applied in the datasets was 5%.
160 There was also elitism between generations, *i.e.*, 20% of the best solutions
161 survive to the next generation.

162 3. Experimental Results

163 The results were found using R programming language (R Core Team,
164 2013)². To compare this paper’s proposed algorithm we applied the exact
165 same train and test samples using several well-known learning algorithms
166 for regression, namely: Linear Regression, Support Vector Machines (Meyer
167 et al., 2012), Regression Trees (Therneau et al., 2013), Conditional Infer-
168 ence Trees (Hothorn et al., 2006; Strobl et al., 2007, 2008) and Random
169 Forests (Liaw and Wiener, 2002)

170 In order to train and test the performance of GAPOLY several main-
171 stream datasets were used. For each dataset, we selected 70% for training
172 purposes and the remaining observations to make a test set in order to com-
173 pute the estimated error. To achieve more robust results, each dataset were
174 processed 25 times, each one with different samples for the train and test sets.
175 For the datasets with attribute values of different magnitudes, a preliminary

²The datasets and the R code used to produce the results and plots in this paper are available online at <https://github.com/jpneto/GenAlgPoly>.

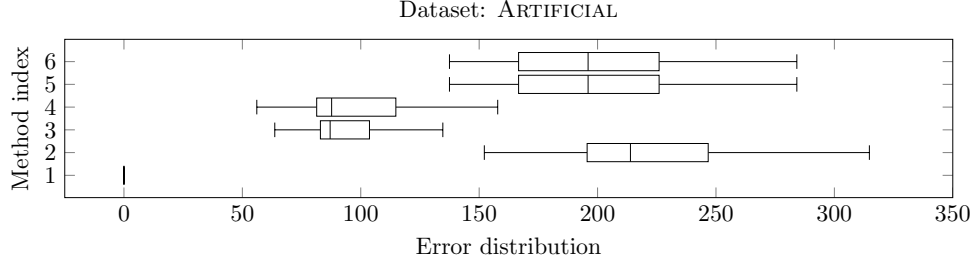


Figure 2: Results for Artificial dataset. As shown, GAPOLY was able to find the exact polynomial structure that generates the dataset, reducing the test set prediction error to zero. The regression methods depicted are: 1. GAPOLY, 2. Random Forests, 3. Linear Regression, 4. SVM, 5. Regression Trees and 6. Conditional Inference Trees

176 scaling was executed. The results below are box plots for the test set error
 177 predictions over these different runs.

178 ARTIFICIAL this is an artificial dataset with four numeric features, x_1, \dots, x_4 ,
 179 where x_1, x_3 are outcomes from Poisson random variables, and x_2, x_4
 180 from Normal random variables. The dependent variable y is given by
 181 expression $x_2x_4^2 + x_1^2x_3 + 5$. The dataset includes $n = 50$ observations.
 182 This dataset was used in order to verify if GAPOLY was able to find
 183 the polynomial relation, which the algorithm did (cf. figure 2). The
 184 genetic algorithm run with a population of $n = 100$ solutions with
 185 50 iterations for each run.

186 In the next datasets, the population of the genetic algorithm had
 187 size $n = 250$ with 100 iterations for each run.

188 HOUSING: This data set concerns the task of predicting housing values
 189 in areas of Boston. There are $m = 13$ continuous attributes and the

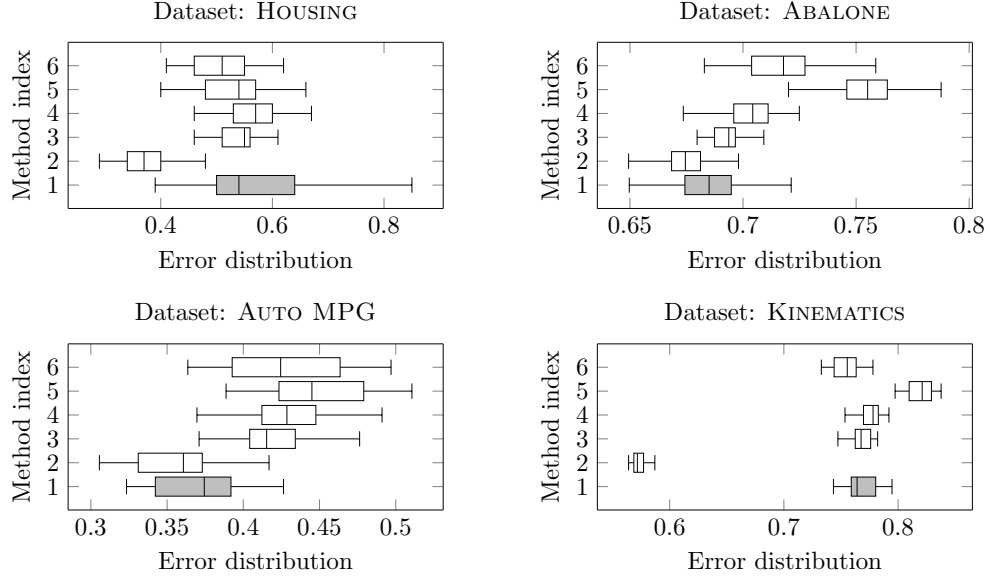


Figure 3: Summary results for different regression methods on the Housing, Abalone, Auto MPG and Kinematics datasets. Random Forests, an ensemble method (“2” in the plots), has the better results. However GAPOLY achieves similar errors in two of the datasets. For the remaining two GAPOLY produces similar errors to the other classic regression methods. The regression methods depicted in these figures are: 1. GAPOLY, 2. Random Forests, 3. Linear Regression, 4. SVM, 5. Regression Trees and 6. Conditional Inference Trees

190 dependent variable is the median value of owner-occupied homes in
 191 \$1000’s. There are $n = 506$ observations.

192 Just as an example of the model the GAPOLY algorithm outputs: in
 193 this dataset the best polynomial,

$$y = -0.12x_6x_9^4 + 0.78x_6 - 0.4x_9^2x_{13} + 0.17x_{13}^2 - 0.044$$

194 where the attributes mean: x_6 , RM average number of rooms per

dwelling; x_9 , RAD index of accessibility to radial highways; x_{13} , Lower status of the population.

For comparison, if we access the mean decrease in accuracy found by Random Forests, the most important attributes are — in decreasing order — x_{13}, x_6, x_5 (but x_5 is already considered 4 times less important than x_6). Both algorithms agree in two of their three most important attributes.

ABALONE This dataset can be used to predict the age of a abalone shell using the given $m = 8$ numeric attributes concerning several physical measurements. There are $n = 4177$ observations.

AUTO MPG This dataset is used to predict fuel consumption in miles per gallon, based on two discrete and five continuous attributes ($m = 7$). There are $n = 398$ observations.

KINEMATICS This dataset is concerned with the realistic simulation of the forward kinematics of an 8 link robot arm. The task is to predict the distance of the end-effector from a target using $m = 8$ continuous attributes. There are $n = 8192$ observations.

3.1. Convergence speed

The GA quickly proceeds in the first 50 to 100 generations to reasonable error rates. Then, it proceeds slower achieving best solutions with marginal error reduction. Since the entire learning process takes some time, in the current R implementation, placing a limit between 50 to 100 generations already achieves good results, relative to higher iteration values. Figure 4

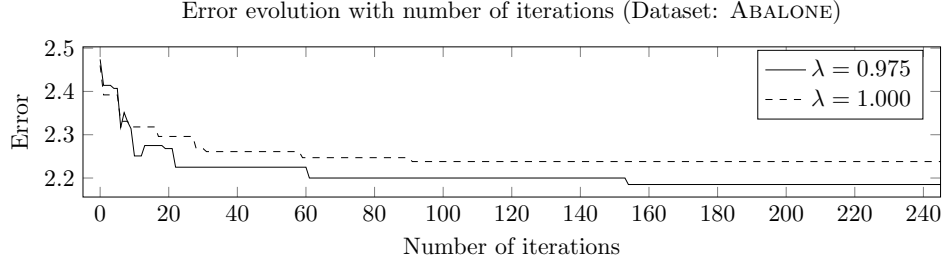


Figure 4: Error progress for Abalone dataset during a single execution of the genetic algorithm. The figure shows the fitness evolution for two different regularization values. The population for both consisted of 200 polynomials. The error values seem to stabilize around iteration 250.

218 shows a typical error evolution for the dataset Abalone given two different
 219 values for λ .

220 4. Conclusion and Future Work

221 Of all the non-ensemble regression methods considered, the proposed
 222 method is the one that shows, with one exception, the lowest error. Only
 223 Random Forest outperforms GAPOLY systematically (it also outperforms
 224 all the other regression algorithms). The single non-ensemble exception —
 225 besides the artificial dataset that uses a straightforward polynomial relation
 226 — is the Auto MPG dataset where GAPOLY has comparable results. This
 227 is evidence that applying standard genetic operators for polynomial model
 228 searching is a viable tool for regression purposes.

229 For complexity considerations GAPOLY demands some processing time.
 230 On a current quad-core computer, processing the Kinematics dataset (with
 231 8k observations) takes approximately 5 minutes. The processing time can

probably be speeded by one to two orders in magnitude if the process is implemented in a low level programming language like C++. However, speed optimization was not the focus of this article.

A cross-validation procedure can be implemented to refine the appropriate parameter values to achieve better errors. Namely, the regularization parameter, λ , can be tested with several values, instead of being fixed at 0.8. Other parameters like mutation chance or the amount of elitism could also be tested. However, these type of tests need a low-level, fast implementation of GAPOLY.

Acknowledgements

The authors are grateful to the Fundação para a Ciência e Tecnologia (FCT) and the R&D laboratory LabMAg for the financial support given to this work, under the strategic project PEST-OE/EEI/UI0434/2011.

The datasets used herein were selected from Luís Torgo's data repository, <http://www.dcc.fc.up.pt/~ltorgo/Regression/DataSets.html>. Most come originally from UCI ML repository, <http://archive.ics.uci.edu/ml/>.

The authors wish to thank professor André Falcão for motivation and useful discussions around the article.

Barricelli, N.A., 1962. Numerical testing of evolution theories. part i: Theoretical introduction and basic tests. *Acta Biotheoretica* 16, 69–98.

Bengio, Y., 2009. Learning deep architectures for AI. *Foundations and trends in Machine Learning* 2, 1–127.

- 254 Bengio, Y., Courville, A., Vincent, P., 2013. Representation learning: A
255 review and new perspectives .
- 256 Cetisli, B., Kalkan, H., 2011. Polynomial curve fitting with varying real
257 powers. *Electronics and Electrical Engineering* 112, 117–122.
- 258 Chan, K.Y., Dillon, T.S., Kwong, C.K., 2011. Modeling of a liquid epoxy
259 molding process using a particle swarm optimization-based fuzzy regression
260 approach. *Industrial Informatics, IEEE Transactions on* 7, 148–158.
- 261 Chan, K.Y., Kwong, C., Dillon, T.S., 2012. Development of product design
262 models using fuzzy regression based genetic programming, in: *Computational Intelligence Techniques for New Product Design*. Springer, pp.
263 111–128.
264
- 265 Cséfalvayová, L., Pelikan, M., Kralj Cigić, I., Kolar, J., Strlič, M., 2010.
266 Use of genetic algorithms with multivariate regression for determination of
267 gelatine in historic papers based on FT-IR and NIR spectral data. *Talanta*
268 82, 1784–1790.
- 269 Davidson, J., Savic, D., Walters, G., 1999. Method for the identification of
270 explicit polynomial formulae for the friction in turbulent pipe flow. *Journal*
271 *of Hydroinformatics* 1, 115–126.
- 272 Gálvez, A., Iglesias, A., Puig-Pey, J., 2012. Iterative two-step genetic-
273 algorithm-based method for efficient polynomial b-spline surface recon-
274 struction. *Information Sciences* 182, 56–76.
- 275 García Nieto, P.J., Alonso Fernández, J., de Cos Juez, F., Sánchez Lasheras,
276 F., Díaz Muñoz, C., 2013. Hybrid modelling based on support vector

277 regression with genetic algorithms in forecasting the cyanotoxins presence
 278 in the trasona reservoir (northern spain). Environmental research .

279 Hofwing, M., Strömberg, N., Tapankov, M., 2011. Optimal polynomial
 280 regression models by using a genetic algorithm, in: Proceedings of the
 281 Second International Conference on Soft Computing Technology in Civil,
 282 Structural and Environmental Engineering Conference,(Crete, Greece),
 283 2011009.

284 Holland, J.H., 1975. Adaptation in natural and artificial systems: An intro-
 285 ductory analysis with applications to biology, control, and artificial intel-
 286 ligence. U Michigan Press.

287 Hothorn, T., Hornik, K., Zeileis, A., 2006. Unbiased recursive partitioning: A
 288 conditional inference framework. Journal of Computational and Graphical
 289 Statistics 15, 651–674.

290 Koza, J.R., 1992. Genetic Programming: vol. 1, On the programming of
 291 computers by means of natural selection. volume 1. MIT press.

292 Liaw, A., Wiener, M., 2002. Classification and regression by randomforest.
 293 R News 2, 18–22. URL: <http://CRAN.R-project.org/doc/Rnews/>.

294 Maertens, K., De Baerdemaeker, J., Babuška, R., 2006. Genetic polynomial
 295 regression as input selection algorithm for non-linear identification. Soft
 296 Computing 10, 785–795.

297 Mendes, L., Carvalho, P.d., 2005. Adaptive polynomial regression for colori-
 298 metric scanner calibration using genetic algorithms, in: Intelligent Signal
 299 Processing, 2005 IEEE International Workshop on, IEEE. pp. 22–27.

300 Meyer, D., Dimitriadou, E., Hornik, K., Weingessel, A., Leisch, F., 2012.
 301 e1071: Misc functions of the department of statistics (e1071), tu wien URL:
 302 <http://CRAN.R-project.org/package=e1071>. r package version 1.6-1.

303 R Core Team, 2013. R: A language and environment for statistical computing
 304 URL: <http://www.R-project.org/>.

305 Sánchez, L., Otero, J., Couso, I., 2009. Obtaining linguistic fuzzy rule-
 306 based regression models from imprecise data with multiobjective genetic
 307 algorithms. *Soft Computing* 13, 467–479.

308 Strobl, C., Boulesteix, A.L., Kneib, T., Augustin, T., Zeileis, A., 2008. Con-
 309 ditional variable importance for random forests. *BMC Bioinformatics* 9.
 310 URL: <http://www.biomedcentral.com/1471-2105/9/307>.

311 Strobl, C., Boulesteix, A.L., Zeileis, A., Hothorn, T., 2007. Bias in ran-
 312 dom forest variable importance measures: Illustrations, sources and a so-
 313 lution. *BMC Bioinformatics* 8. URL: [http://www.biomedcentral.com/](http://www.biomedcentral.com/1471-2105/8/25)
 314 [1471-2105/8/25](http://www.biomedcentral.com/1471-2105/8/25).

315 Tarlow, D., Sutskever, I., Zemel, R.S., 2013. Stochastic k-neighborhood
 316 selection for supervised and unsupervised learning. *Journal of Machine*
 317 *Learning Research* .

318 Therneau, T., Atkinson, B., Ripley, B., 2013. rpart: Recursive partitioning
 319 R package version 4.1-1.

320 Willighagen, E., 2012. genalg: R based genetic algorithm.

- 321 Wu, C.H., Tzeng, G.H., Lin, R.H., 2009. A novel hybrid genetic algorithm for
322 kernel function and parameter optimization in support vector regression.
323 Expert Systems with Applications 36, 4725–4735.
- 324 Yu, T.L., Lin, W.K., 2008. Optimal sampling of genetic algorithms on poly-
325 nomial regression, in: Proceedings of the 10th annual conference on Ge-
326 netic and evolutionary computation, ACM. pp. 1089–1096.